

Programming Assignment 4

Query Me This

Time due: 9:00 PM Wednesday, June 9

For this project, you will build a simple database in memory and process database queries.

Introduction

A simple database *table* is a collection of *records*. Each record has the same number of *fields*; we'll call that the number of *columns* in the table. Each of the columns has a name. For example, we may have a 4-column table whose columns are named `customer`, `product`, `price`, and `location`. Here are some records we might have in this table:

customer	product	price	location
Patel	12345	42.54	Westwood
O'Reilly	34567	4.99	Westwood
Hoang	12345	30.46	Santa Monica
Patel	67890	142.75	Hollywood

One database operation of interest is the one that finds all records in a table for which one particular field equals a certain value. We make the simplifying assumption for this project that for any particular table, there's only one field for which we want such lookups to be especially fast. We call this field the *key field*. Let's suppose for this example that we designated the `customer` field as the key field. Then we want the operation "Find all records where the customer is Patel" to quickly produce the records:

```
Patel 12345 42.54 Westwood
Patel 67890 142.75 Hollywood
...
```

Notice that we allow more than one record to have the same key.

Another operation is a more general query capability. We may want to select all records for which, for example, the price is less than 40. This would produce

```
O'Reilly 34567 4.99 Westwood
Hoang 12345 30.46 Santa Monica
...
```

Your Task

You will implement a class `Table` that will provide the functionality described above.

class Table

The `Table` class is responsible for holding table records and responding to queries about them. Records **must** be stored in a data structure suitable for fast insertions and lookups by key: This means a tree or hash table. If you choose a tree, you do not have to worry about balancing it; if you choose a hash table, you do not have to worry about changing the number of buckets dynamically. Also, if you choose a hash table, it must not have more than 1000 buckets (to make it easier for us when we test your program on a table with thousands of records). For this project, you must implement the tree or hash table data structure yourself. You must **not** use any STL associative container: no `map`, `multimap`, `unordered_map`, `unordered_multimap`, `set`, `multiset`, `unordered_set`, or `unordered_multiset`. You *are* allowed to use `vector`, `list`, `stack`, and `queue`.

```
class Table
{
public:
    Table(std::string keyColumn, const std::vector<std::string>& columns);
    ~Table();
    bool good() const;
    bool insert(const std::string& recordString);
    void find(std::string key, std::vector<std::vector<std::string>>& records) const;
    int select(std::string query, std::vector<std::vector<std::string>>& records) const;

    // We prevent a Table object from being copied or assigned by
    // making the copy constructor and assignment operator unavailable.
    Table(const Table&) = delete;
    Table& operator=(const Table&) = delete;
private:
    ...
};
```

You must **not** make any deletions, additions, or changes to the *public* interface of this class. You can and will, of course, add private data members and perhaps private member functions to help you with your implementations of the member functions. You must **not** declare any public data members, nor use any global variables whose values may change during execution (so global *constants* are OK). You may add additional functions that are not members of any class.

Table(std::string keyColumn, const std::vector<std::string>& columns)

Construct an empty `Table` whose column names are the elements of the vector second parameter. The first parameter is the name of the key field. The parameters might be such that you could not construct a valid table; the columns vector might be empty or contain empty or duplicate strings, or the `keyColumn` might not be any of the column names. The right way to handle such a problem in modern C++ is to throw an exception, but we haven't discussed exceptions. Instead, for this project you must set the state of the object so that it can be tested using the `good` member function. For example, a user might write:

```
Table t(constructor argument);
if ( ! t.good() )
    give up, since the constructor failed
```

This isn't a great solution, since a careless user could forget to test for the failure and try to use an invalid `Table` object, but it's the best we can do without exceptions. To mitigate the problem, each member function should check that it's been called on a valid object, and return harmlessly, if possible, if it hasn't.

bool good() const

Return true if the table was successfully constructed, and false otherwise.

~Table()

The destructor does what is necessary to ensure that a `Table` releases any resources it holds when its lifetime ends.

bool insert(const std::string& recordString);

Insert a record into the database. The parameter is a string with the fields of the record in a format described below. If the string has the correct number of fields for the table, insert a record with those fields into the table and return true; otherwise, leave the table unchanged and return false. Here is an example:

```
// Since C++11, vectors can be constructed from initializer lists
vector<string> cols = {
    "customer", "product", "price", "location"
};
Table t("customer", cols);
assert(t.good());
assert(t.insert("Patel 12345 42.54 Westwood"));
assert(t.insert("O'Reilly 34567 4.99 Westwood "));
assert(t.insert(" Hoang 12345 30.46 'Santa Monica' "));
assert(t.insert("Patel\t67890\t142.75  \t \t\t \tHollywood"));
assert( ! t.insert("Figueroa 54321 59.95"));
```

A *delimiter* is a blank, tab, carriage return, or newline character. A field in the `recordString` is represented by either

- a sequence of one or more non-delimiter characters, not starting with a single quote; or
- a single quote followed by zero or more of the following units followed by a single quote:
 - a character other than a single quote; or
 - two consecutive single quotes (representing one single quote in the actual field value)

Here are some examples:

recordString textrepresents the field value
Patel	Patel
'Santa Monica'	Santa Monica
chillin'	chillin'
O'Reilly	O'Reilly
'chillin'''	chillin'
'O'Reilly'	O'Reilly
'Baba O'Riley – Who's Next'	Baba O'Riley – Who's Next
''	<i>the empty string</i>

In the `recordString`, the representations of the fields are separated by one or more delimiters, and the `recordString` starts or ends with zero or more delimiters. One exception: If a field representation surrounded by single quotes is followed by a field representation not surrounded by single quotes, they need not be separated by delimiters, so the two field representations `'da gama' 98765` could also appear in the string as `'da gama'98765`.

void find(std::string key, std::vector<std::vector<std::string>>& records) const;

Replace the value in the second parameter with a vector containing as many elements as there are records in the table whose key field is equal to the value of the first parameter. Each of those elements is one of the matching records, represented as a vector whose elements are the fields of that record. The records placed in the vector don't have to be in any particular order. For example, for the table `t` built in the example shown for the `insert` function, the assertions in the following should succeed:

```
vector<vector<string>> v;
t.find("Patel", v);
assert(v.size() == 2);
vector<vector<string>> expected = {
    { "Patel", "12345", "42.54", "Westwood" },
    { "Patel", "67890", "142.75", "Hollywood" }
};
assert((v[0] == expected[0]  &&  v[1] == expected[1])  ||
       (v[0] == expected[1]  &&  v[1] == expected[0]) );

int select(std::string query, std::vector<std::vector<std::string>>& records) const;
```

The specification of this function will appear Saturday night.

Other Requirements

Your code will go into two files: `Table.h`, which will contain the declaration of the `Table` class and must have an appropriate include guard, and `Table.cpp`, which will contain the implementation of the `Table` class and any additional functions you might write to help you with the implementation. If you want to have a main routine to test your code, put it in a separate `.cpp` file that you will not turn in.

You must have an implementation for every function specified by this assignment. If you can't get a function implemented correctly, it must at least compile successfully. For example, if you don't have time to correctly implement `Table::select`, say, here is an implementation that meets this requirement in that it at least compiles correctly:

```
int select(string query, vector<vector<string>> records) const
{
    return 0; // not always correct, but at least this compiles
}
```

None of the member functions you write may cause anything to be read from `cin` or written to `cout`. If you want to print things out for debugging purposes, write to `cerr` instead of `cout`. When we test your program, we will cause everything written to `cerr` to be discarded; we will never see that output, so you may leave those debugging output statements in your program if you wish.

Your code must compile successfully under both `g32` and either `Visual C++` or `clang++`. If your code is linked to a file containing

```
#include "Table.h"
#include <iostream>
#include <cassert>
using namespace std;

int main()
{
    vector<string> cols = { "N", "Z" };
    Table t("Z", cols);
    assert(t.good());
    assert(t.insert("UCLA 90095"));
    assert(t.insert("Caltech 91125"));
    vector<vector<string>> v;
    t.find("90095", v);
    assert(v.size() == 1);
    assert(v[0][0] == "UCLA"  &&  v[0][1] == "90095");
    t.select("", v); // This will be changed Saturday
    cout << "DONE" << endl;
}
```

the linking must succeed. When the resulting executable is run, execution must at least reach the second assert. (Of course, if you've correctly implemented the functions, the program will run to completion.)

- During execution, if a client uses the classes as defined, your program must not perform any undefined actions, such as dereferencing a null or uninitialized pointer, or accessing a vector element out of bounds.

Implementation Hints

To help you extract fields from a `recordString`, we've written this [StringParser class](#). Read it to figure out what it does, use it, and save a lot of work.

Every STL container has a member function named `clear` that removes all elements from the container, leaving it empty.

The header `<functional>` defines a hash class template that you can use to hash strings:

```
string s("Hello there");
unsigned int h = std::hash<std::string>()(s);
```

The hash values produced by this function range from 0 to about 4 billion.

Turn it in

By Tuesday, June 8, there will be a link on the class webpage that will enable you to turn in your source files. You will turn in a zip file containing these two source code files and nothing more:

- `Table.h`
- `Table.cpp`

You do not have to turn in a report file.